

Workshop on **Hardware and Algorithms for Learning On-a-chip (HALO) 2019**

The Westin Westminster, Westminster, CO, Thursday, November 7, 2019

Time	Schedule
8:15am – 8:30am	Introduction and opening remarks
8:30am – 9:20am	Keynote talk 1: High Performance Mapping for Deep Learning Algorithms Allen Rush (AMD)
9:20am – 10:10am	Keynote talk 2: Machine Learning on Social Network Platforms Hsien-Hsin Lee (Facebook)
10:10am – 10:30am	Coffee Break
Session 1: Hardware Acceleration of Machine Learning	
10:30am – 10:50am	Neural Networks Accelerator Design from the User Perspective Yu Wang (Tsinghua University)
10:50am – 11:10am	Algorithm/Hardware Co-design for Energy/Area Efficient In-Memory Neural Network Computing Jae-Joon Kim (POSTECH)
11:10am – 11:30am	Bringing Powerful Machine-learning Systems to Daily-life Devices via Algorithm-Hardware Co-design Yingyan Lin (Rice University)
11:30am – 11:50am	Toward Next-Generation Acceleration for AI: A Heterogeneous Computing Approach Jaewoong Sim (Intel)
11:50am – 12:10pm	Bayesian Tensorized Neural Networks with Automatic Rank Determination Zheng Zhang (UCSB)
12:10pm – 1:10pm	Lunch
1:10pm – 2:00pm	Keynote talk 3: On-Device AI for Augmented Reality (AR) Systems Vikas Chandra (Facebook)
Session 2: Intelligent Mobile Applications	
2:00pm – 2:20pm	A Product Engine for Energy-Efficient Execution of Binary Neural Networks Using Resistive Memories Pierre-Emmanuel Gaillardon (University of Utah)
2:20pm – 2:40pm	Evolutionary Optimization for Neuromorphic Systems Catherine Schuman (Oak Ridge National Lab)
2:40pm – 3:00pm	Micro AI: When Intelligence Moves to the Low Power Sensors Tinoosh Mohsenin (UMBC)
3:00pm – 3:20pm	Beyond Energy-Efficiency: Enabling Fault-Aware Learning On-a-chip Siddharth Garg (New York University)
3:20pm – 3:40pm	Coffee Break

Poster Session

<p>3:40pm – 4:10pm Poster presentations (2 mins each)</p>	<p>P1. How to Obtain and Run Light and Efficient Deep Learning Networks Fan Chen (Duke University)</p> <p>P2. Achieving Super-Linear Speedup across Multi-FPGA for Real-Time DNN Inference Weiwen Jiang (University of Notre Dame)</p> <p>P3. On Neural Architecture Search for Resource-Constrained Hardware Platforms Qing Lu (University of Notre Dame)</p> <p>P4. CIMAT: A Transpose SRAM-based Compute-In-Memory Architecture for Deep Neural Network On-Chip Training Hongwu Jiang (Georgia Institute of Technology)</p> <p>P5. HR3AM: A Heat Resilient Design for RRAM Based Neuromorphic Computing Xiao Liu (University of California San Diego)</p> <p>P6. ACG-Engine: An Inference Accelerator for Content Generative Neural Networks Haobo Xu (Institute of Computing Technology, Chinese Academy of Sciences)</p> <p>P7. Mixed Precision Neural Architecture Search for Energy Efficient Deep Learning Zhixuan Jiang (The University of Texas at Austin)</p> <p>P8. Enhanced Error-Correcting DNN Classifier Towards Robust Machine Learning On-a-chip Tao Liu (Lehigh University)</p> <p>P9. PCONV: A Desirable Sparsity Dimension for Real-time Execution -- From Algorithm to Framework Xiaolong Ma (Northeastern University)</p> <p>P10. 2.5ms MobileNet-V2 Execution on Mobile Phone -- A Compiler-Assisted Block Pruning Framework Zhengang Li (Northeastern University)</p> <p>P11. INA: Incremental Network Approximation Algorithm for Limited Precision Deep Neural Networks Zheyu Liu (Tsinghua University)</p> <p>P12. Approximating Backpropagation for a Biologically Plausible Local Learning Rule in Spiking Neural Networks Haowen Fang (Syracuse University)</p> <p>P13. Leveraging Model Diversity for High QoS Deep Learning Inference in the Clouds Jeff (Jun) Zhang (New York University)</p>
<p>4:10pm – 5:00pm</p>	<p>Poster discussions</p>